

Zentry Service Architecture

Document Version: 1.0

Date: April 30, 2025

Author: Team Zentry

Table of Contents

1. [Executive Summary](#)
2. [System Overview](#)
3. [Service Architecture](#)
4. [Data Flow](#)
5. [Component Interactions](#)
6. [Scaling Considerations](#)
7. [Deployment Strategy](#)
8. [Monitoring and Observability](#)
9. [Appendix](#)

Executive Summary

Zentry is a sleep tracking social application that allows users to log their sleep patterns and follow other users to view their sleep entries. The application employs a distributed microservice architecture with event-driven processing to handle data consistency across services while maintaining scalability and reliability.

This document details the service architecture, explaining how the various components work together to deliver a robust, scalable application that can handle both regular users and highly-followed “celebrity” users efficiently.

System Overview

Purpose and Scope

Zentry provides functionality for:

- Recording and managing sleep entries
- Following/unfollowing other users
- Viewing a personalized feed of sleep entries from followed users
- Analyzing sleep patterns and statistics

Technical Stack

- **Backend Framework:** Ruby on Rails 8.0 (API-only mode)
- **Primary Database:** PostgreSQL (relational data)
- **Search & Analytics:** Elasticsearch 7.13 (feed and search functionality)
- **Message Broker:** Kafka (event processing)
- **Containerization:** Docker
- **Deployment:** Kamal

Design Principles

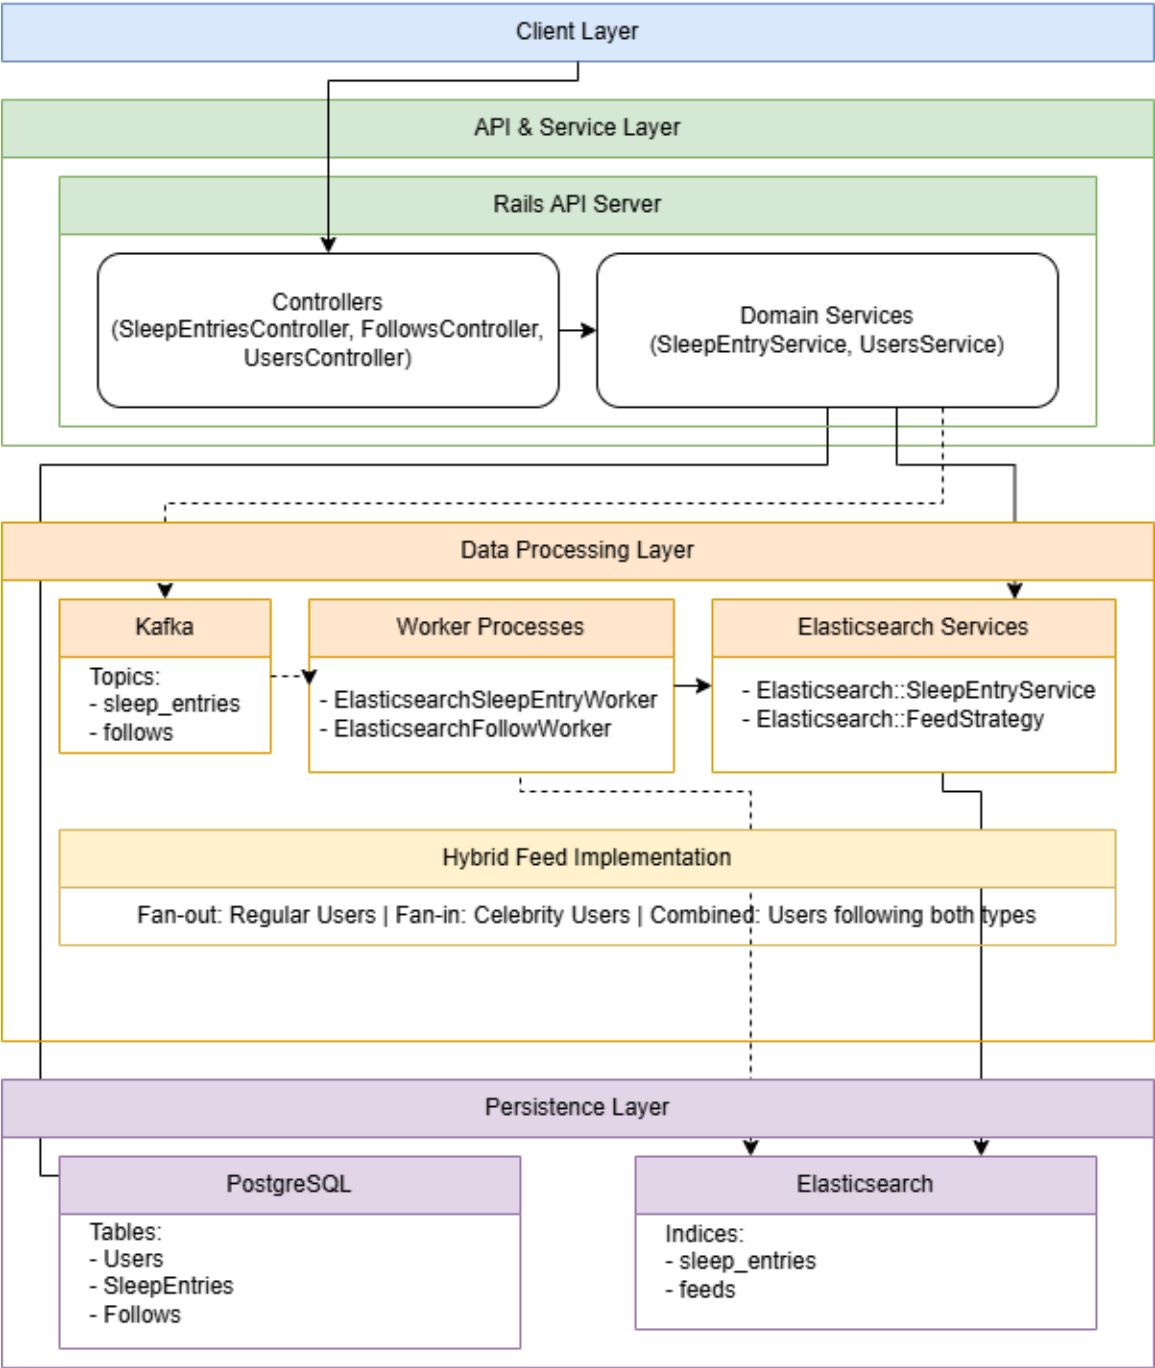
The architecture follows these core principles: - Event-driven processing for asynchronous operations - Hybrid fan-out/fan-in approach for social feed implementation - Separation of read and write paths to optimize performance - Domain-driven design with service boundaries

Service Architecture

Core Components

Service Architecture Diagram

Zentry Service Architecture



Zentry uses a hybrid fan-out/fan-in architecture to balance read and write performance

Legend
—— Synchronous Call
- - - Asynchronous Event
→ Data Flow Direction

The system consists of the following major components:

1. **API Layer:** Rails controllers handling HTTP requests
2. **Domain Services:** Business logic implementation
3. **Data Access Layer:** Models and repositories
4. **Event Processing System:** Kafka producers and consumers
5. **Search & Feed Infrastructure:** Elasticsearch services and indices
6. **Background Workers:** Service workers for asynchronous processing

Component Details

API Layer

The API layer provides RESTful endpoints for all client interactions:

- **SleepEntriesController:** CRUD operations for sleep entries
- **FollowsController:** Managing follow relationships
- **UsersController:** User profile and authentication operations

Domain Services

Services encapsulate the core business logic:

- **SleepEntryService:** Management of sleep entry creation, updates and retrieval
- **UserService:** User relationship management and feed generation
- **Elasticsearch::SleepEntryService:** Search and feed assembly from Elasticsearch

Data Access Layer

The persistence layer uses both PostgreSQL and Elasticsearch:

- **PostgreSQL Models:**
 - User: User account information
 - SleepEntry: Sleep record data
 - Follow: Social graph relationships
- **Elasticsearch Indices:**
 - sleep_entries: Primary storage for searchable sleep entries
 - feeds: Denormalized feed entries for regular users

Event Processing System

The event system handles data consistency across services:

- **Kafka Topics:**
 - sleep_entries: Events related to sleep entry operations
 - follows: Events related to social graph changes
- **Producers:**

- `Kafka::Producer`: Generic producer for publishing events
- **Consumers:**
 - `ElasticsearchSleepEntryWorker`: Consumes sleep entry events
 - `ElasticsearchFollowWorker`: Consumes follow relationship events

Data Flow

Create Sleep Entry Flow

1. Client makes a POST request to `/sleep_entries`
2. `SleepEntriesController` validates input and calls `SleepEntryService.create`
3. `SleepEntryService` creates a record in PostgreSQL
4. After successful creation, a `sleep_entry_created` event is published to Kafka
5. `ElasticsearchSleepEntryWorker` consumes the event and:
 - Indexes the entry in the `sleep_entries` Elasticsearch index
 - For non-celebrity users, copies the entry to each follower's feed in the feeds index (fan-out)

Follow User Flow

1. Client makes a POST request to `/follows`
2. `FollowsController` validates input and calls `UserService.create_follow`
3. `UserService` creates a follow relationship in PostgreSQL
4. After successful creation, a `follow_created` event is published to Kafka
5. `ElasticsearchFollowWorker` consumes the event and:
 - For non-celebrity users, copies all of their sleep entries to the follower's feed (fan-out)
 - For celebrity users, no immediate action (fan-in strategy)

View Feed Flow

1. Client makes a GET request to `/feed`
2. `SleepEntriesController` calls `SleepEntryService.get_followers_feed`
3. `SleepEntryService` delegates to `Elasticsearch::SleepEntryService.feed_for_user`
4. `Elasticsearch::SleepEntryService` identifies the user's mix of followed users and:
 - For regular followed users: queries pre-computed entries in the feeds index (fan-out)
 - For celebrity followed users: queries their entries directly from `sleep_entries` index (fan-in)
 - Combines both result sets, sorts, and paginates for final response

Component Interactions

Hybrid Feed Implementation

The system implements a hybrid approach for feed generation:

1. **Fan-out pattern:**
 - Used for regular users with moderate follower counts
 - When a user posts a new sleep entry, it's copied to each follower's feed
 - Improves read performance at the cost of write amplification
 - Implemented in `ElasticsearchSleepEntryWorker.index_sleep_entry`
2. **Fan-in pattern:**
 - Used for “celebrity” users with large follower counts
 - Sleep entries are stored only once in the `sleep_entries` index
 - When a follower requests their feed, celebrity entries are fetched on-demand
 - Optimizes write performance while slightly increasing read complexity
 - Implemented in `Elasticsearch::SleepEntryService.fan_in_feed_query`
3. **Combined approach:**
 - For users following both regular and celebrity accounts
 - Results from both approaches are merged and sorted at query time
 - Ensures consistent behavior regardless of the followed users' profile
 - Implemented in `Elasticsearch::SleepEntryService.combined_feed_query`

Asynchronous Event Processing

The system uses Kafka for reliable event handling:

1. **Transaction with Kafka:**
 - Database operations and Kafka publishing are wrapped in transactions
 - Ensures consistency between database state and events
 - Implemented in service patterns like `UserService.create_follow`
2. **Idempotent Workers:**
 - Workers are designed to handle duplicate messages safely
 - Each operation is idempotent, allowing for at-least-once delivery semantics
 - Retry mechanisms ensure eventual consistency

Scaling Considerations

Horizontal Scaling

The architecture supports horizontal scaling through:

1. **Stateless API Servers:**
 - Can be scaled independently based on request volume

2. **Separate Worker Processes:**
 - Kafka consumers can be deployed separately from API servers
 - Consumer groups ensure work distribution across worker instances
3. **Elasticsearch Cluster:**
 - Supports sharding and replication for performance and reliability
 - Can be scaled independently of other components

Performance Optimizations

1. **Celebrity Detection:**
 - FeedStrategy service determines optimal feed strategy based on follower count
 - Prevents write amplification for highly followed accounts
2. **Query Optimizations:**
 - Sorting and pagination handled efficiently
 - Date range filtering at the query level
3. **Bulk Operations:**
 - Bulk indexing for efficiency when processing large data sets

Deployment Strategy

Infrastructure Components

The production deployment consists of:

1. **Application Servers:**
 - Multiple instances of the Rails application
 - Deployed and managed using Kamal
2. **Data Stores:**
 - PostgreSQL database cluster with replication
 - Elasticsearch cluster with proper sharding
3. **Kafka Cluster:**
 - Multiple brokers for reliability
 - Zookeeper ensemble for coordination

Deployment Process

WIP

Monitoring and Observability

Logging Strategy

The application employs structured logging:

1. **Application Logs:**
 - Rails application logs

- Worker-specific logs
- 2. **Service Logs:**
 - Elasticsearch operational logs
 - Kafka broker and consumer group logs

Performance Monitoring

Monitoring focuses on:

1. **API Performance:**
 - Request latency
 - Error rates
2. **Worker Performance:**
 - Message processing rates
 - Consumer lag
3. **Search Performance:**
 - Query performance
 - Index health metrics

Appendix

Application Initialization Flow

1. Rails server starts
2. `elasticsearch_worker.rb` initializer starts worker threads:
 - `ElasticsearchSleepEntryWorker` for sleep entry events
 - `ElasticsearchFollowWorker` for follow relationship events
3. Workers connect to Kafka and begin consuming messages

Database Schema

Key tables and relationships:

- `users`: User profiles and authentication
- `sleep_entries`: Sleep record data linked to users
- `follows`: Social graph with follower/following relationships

Elasticsearch Index Mappings

- **sleep_entries index:**

```
{
  "properties": {
    "id": { "type": "long" },
    "user_id": { "type": "long" },
    "sleep_entry_id": { "type": "long" },
    "sleep_duration": { "type": "long" },
    "sleep_start_at": { "type": "date" },
    "created_at": { "type": "date" },
  }
}
```



```
    "updated_at": { "type": "date" }  
  }  
}
```

- **feeds index:**

```
{  
  "properties": {  
    "user_id": { "type": "long" },  
    "author_id": { "type": "long" },  
    "sleep_entry_id": { "type": "long" },  
    "sleep_duration": { "type": "long" },  
    "sleep_start_at": { "type": "date" },  
    "created_at": { "type": "date" },  
    "updated_at": { "type": "date" }  
  }  
}
```